

CSE 451: Operating Systems

Section 1

Intro, C programming, project 0

Introduction

- * My name is Sean and I am a senior in Computer Science and Biochemistry
- * I worked in the Multimedia Platform Team on Capture Performance at Microsoft this summer
- * My background is mainly in Experimental and Computational Biochemistry
- * I enjoy C/C++ programming, especially template errors, and systems
 - * Favorite classes: Databases, OS, Comp Bio
 - * I have been a TA for CSE 341 and for Coursera Courses Introduction to Data Science and Programming Languages
 - * Currently TA for a Coursera course this quarter with Dan Grossman, so join if you have time! It's essentially 341 online
<https://www.coursera.org/course/proglang>
- * My office hours are Mondays, 10:30-11:20, Fridays, 1:30-2:20 or by appointment. Or you can tackle me in the halls wherever you see me.
- * Contact: discussion board or by email (wujsean@cs)

Why are you here?

- * You need a 400-level elective and this course fit your schedule
- * You have heard that Ed is an exceptional lecturer
- * You want a job / to go to graduate school in CS
- * You want to understand what goes on “under the hood”

Far-reaching implications

- * Concepts and techniques learned in lecture / through projects apply to all other areas of computer science
 - * Data structures
 - * Caching
 - * Concurrency
 - * Virtualization
- * OSes *support* all other areas of computer science

Course tools

- * Assn 0: Any computer with C development tools (002, attu, your *nix box)
- * Assn 1: Use the course VM inside an emulator (VMware, VirtualBox, Qemu etc.) on your computer or a lab computer
- * Can compile on forkbomb.cs.washington.edu (faster)

Course tools

- * We'll be using the GNU C Compiler (gcc) for compiling C code in this course, which is available on every platform except Windows (Cygwin lovers proceed at your own risk)
- * For an editor, use whatever you are most comfortable with; emacs, vim, gedit, and Eclipse are good choices (ed and butterflies also options)

Discussion board

- *The discussion board is an invaluable tool; use it!
- *Jeff (my TA partner in crime) and I both receive email alerts whenever there is a new post. Response time should be very fast.
- *For anything non-personal use the discussion board.

Collaboration

- * If you talk or collaborate with anybody, or access any websites for help, *name them* in your project submission
- * See the course policy for more restrictions
- * Okay: discussing problems and techniques to solve them with other students
- * Not okay: looking at/copying other students' code. Googling solutions. Using code from Wikipedia.
- * We will pass your code through plagiarism detection software (MOSS, Deckard, etc.)

C programming

- * Most modern operating systems are still written in C
- * Why not Java?
 - * Interpreted Java code runs in a virtual machine, so what language is the VM built in?
- * C is precise in terms of
 - * Instructions (semantics are clear)
 - * Timing (can usually estimate number of cycles needed to execute code)
 - * Memory (allocations/de-allocations are explicit)

C language features

- * Pointers

- * Pass-by-value vs. pass-by-reference

- * Structs

- * Typedefs (aliasing)

- * Malloc/free

Pointers

```
int iX = 5;
```

```
int iY = 6;
```

```
int* piX = &iX;    // declare a pointer to x  
                  // with value as the  
                  // address of x
```

```
*piX = iY;         // change value of x to y  
                  // (x == 6)
```

```
piX = &iY;          // change px to point to  
                  // y's memory location
```

```
// For more review, see the CSE 333 lecture  
// and section slides
```

Function pointers

```
int functionate(int x, char c) { ... }  
    // declare and define a function  
int (*pfFoo)(int, char) = NULL;  
    // declare a pointer to a function  
    // that takes an int and a char as  
    // arguments and returns an int  
pfFoo = functionate;  
    // assign pointer to  
functionate()'s  
    // location in memory  
iX = pfFoo(7, 'p');  
    // set iX to the value returned by  
    // functionate(7, 'p')
```

Case study: signal()

```
extern void (*signal(int, void(*) (int))) (int);
```

- * What is going on here?

- * `signal()` is "a function that takes two arguments, an integer and a pointer to a function that takes an integer as an argument and returns nothing, and it (`signal()`) returns a pointer to a function that takes an integer as an argument and returns nothing." (from StackOverflow)

Case study: signal

* We can make this a lot clearer using a typedef:

```
// Declare a signal handler prototype
typedef void (*SigHandler)(int iSignum);
// signal could then be declared as
extern SigHandler signal(
    int iSignum, SigHandler pfHandler);
```

Arrays and pointer arithmetic

- * Array variables can often be treated like pointers, and vice-versa:

```
int aiFoo[2];    // foo acts like a pointer to
                 // the beginning of the array
*(aiFoo + 1) = 5; // the second int in the
                 // array is set to 5
```

- * Don't use pointer arithmetic unless you have a good reason to do so

Passing by value vs. reference

```
int doSomething(int iFoo) {  
    return iFoo + 1;  
}
```

```
void doSomethingElse(int* piFoo) {  
    *piFoo += 1;  
}
```

```
void example(void) {  
    int iX = 5;  
    int iY = doSomething(iX);    // iX==5, iY==6  
    doSomethingElse(&iX);        // iX==6, iY==6  
}
```


Returning addl. information

```
int initialize(int iArg1, int iArg2,  
    int* piErrorCode) {  
    // If initialization fails, set an error  
    // code and return false to indicate  
    // failure.  
    if (...) {  
        *piErrorCode = ...;  
        return EXIT_FAILURE;  
    }  
    // ... Do some other initialization work  
    return EXIT_SUCCESS;  
}
```

Structs

```
// Define a struct referred to as
// "struct s2DPoint"
struct s2DPoint {
    int iX;
    int iY;
}; // Don't forget the trailing ';'!

// Declare a struct on the stack
struct s2DPoint foo;

// Set the two fields of the struct
foo.iX = 1;
foo.iY = 2;
```

Typedefs

```
typedef struct s2DPoint 2DPoint;  
    // Creates an alias "2DPoint" for  
    // "struct s2DPoint"  
  
2DPoint* poBar =  
    (2DPoint*) malloc(  
        sizeof(2DPoint));  
    // Allocates space for a 2DPoint struct  
    // on the heap; poBar points to it  
  
poBar->iX = 2;  
    // "->" operator dereferences the  
    // pointer and accesses the field iX;  
    // equivalent to (*poBar).iX = 2;
```

Memory management

- * Allocate memory on the heap:

```
void* malloc(size_t size);
```

- * Note: malloc may fail!

- * But not necessarily when you would expect...

- * Use `sizeof()` operator to get the size of a type/struct

- * Free memory on the heap:

```
void free(void* ptr);
```

- * Pointer argument comes from previous `malloc()` call

Common C pitfalls (1)

* What's wrong and how can it be fixed?

```
char* city_name(float fLat, float fLong) {  
    char sName[100];  
    ...  
    return sName;  
}
```

Common C pitfalls (1)

- * Problem: returning pointer to local (stack) memory (also: using floats)

- * Solution: allocate on heap

```
char* city_name(double fLat, double fLong) {  
    // Preferably allocate a string of  
    // just the right size  
    char* sName =  
        (char*) malloc(100*sizeof(char));  
    ...  
    return sName;  
}
```

Common C pitfalls (2)

* What's wrong and how can it be fixed?

```
char* sBuf = (char*) malloc(32*sizeof(char));  
strcpy(sBuf, argv[1]);
```

Common C pitfalls (2)

- * Problem: potential buffer overflow

- * Solution:

```
static const int BUFFER_SIZE = 32;  
  
char* sBuf = (char*) malloc(BUFFER_SIZE);  
strncpy(sBuf, argv[1], BUFFER_SIZE);
```

- * Why are buffer overflow bugs dangerous?

Common C pitfalls (3)

* What's wrong and how can it be fixed?

```
char* sBuf = (char*) malloc(BUFFER_SIZE);  
Strncpy(sBuf, sHello, BUFFER_SIZE);  
printf("%s\n", sBuf);
```

```
sBuf = (char*) malloc(2*BUFFER_SIZE);  
strncpy(sBuf, sLongHello, 2*BUFFER_SIZE);  
printf("%s\n", sBuf);
```

```
free(sBuf);
```

Common C pitfalls (3)

* Problem: memory leak

* Solution:

```
char* sBuf = (char*) malloc(BUFFER_SIZE);  
strncpy(sBuf, sHello, BUFFER_SIZE);  
printf("%s\n", sBuf);  
free(sBuf);
```

```
buf = (char*) malloc(2*BUFFER_SIZE);  
...
```

Common C pitfalls (4)

* What's wrong (besides ugliness) and how can it be fixed?

```
char sFoo[2];  
sFoo[0] = 'H';  
sFoo[1] = 'i';  
printf("%s\n", sFoo);
```

Common C pitfalls (4)

* Problem: string is not NULL-terminated

* Solution:

```
char sFoo[3];  
sFoo[0] = 'H';  
sFoo[1] = 'i';  
sFoo[2] = '\\0';  
printf("%s\\n", sFoo);
```

* Easier way: `char* sFoo = "Hi";`

Common C pitfalls (5)

- * Another bug in the previous examples?
 - * Not checking the return value of system calls / library calls!

```
char* sBuf = (char*) malloc(BUFFER_SIZE);
if (sBuf == 0) {
    fprintf(stderr, "error!\n");
    return EXIT_FAILURE;
}
strncpy(sBuf, argv[1], BUFFER_SIZE);
...
```

Project 0

- * Description is on course web page
- * Due Friday October 4, 11:59pm
- * Work individually
 - * Remaining projects are in groups of 2. When you have found a partner, one of you should fill out the survey on Catalyst (forthcoming by email)

Project 0 goals

- * Get re-acquainted with C programming
- * Practice working in C / Linux development environment
- * Create data structures for use in later projects

Valgrind

- * Helps find all sorts of memory problems
 - * Lost pointers (memory leaks), invalid references, double frees
- * Simple to run:
 - * `valgrind ./myprogram`
 - * Look for “definitely lost,” “indirectly lost” and “possibly lost” in the LEAK SUMMARY
- * Manual:
 - * <http://valgrind.org/docs/manual/manual.html>

Project 0 memory leaks

- * Before you can check the queue for memory leaks, you should probably add a queue destroy function:

```
void queue_destroy(queue* q) {  
    queue_link* cur;  
    queue_link* next;  
    if (q != NULL) {  
        cur = q->head;  
        while (cur) {  
            next = cur->next;  
            free(cur);  
            cur = next;  
        }  
        free(q);  
    }  
}
```

Project 0 testing

- * The test files in the skeleton code are incomplete
 - * Make sure to test *every* function in the interface (the .h file)
 - * Make sure to test corner cases
- * Suggestion: write your test cases first

Project 0 tips

- * Part 1: queue
 - * First step: improve the test file
 - * Then, use valgrind and gdb to find the bugs
- * Part 2: hash table
 - * Write a thorough test file
 - * Perform memory management carefully
- * You'll lose points for:
 - * Leaking memory
 - * Not following submission instructions
- * Use the discussion board for questions about the code